

Disclaimer

I wrote this to my best knowledge, however, no guarantees are given whatsoever.

Sources

If not noted differently, the source is the lecture slides and/or the accompanying book.

1 Approximate Retrieval

Nearest-Neighbor Find $\mathbf{x}^* = \operatorname{argmin}_{\mathbf{x} \in X} d(\mathbf{x}, \mathbf{y})$ given S , $\mathbf{y} \in S$, $X \subseteq S$.

Near-Duplicate detection Find all $\mathbf{x}, \mathbf{x}' \in X$ with $d(\mathbf{x}, \mathbf{x}') \leq \epsilon$.

1.1 k -Shingling

Documents (or videos) as set of k -shingles (a. k. a. k -grams). k -shingle is consecutive appearance of k chars/words.

Binary *shingle matrix* $M \in \{0,1\}^{C \times N}$ where $M_{i,j} = 1$ iff document j contains shingle i , N documents, C k -shingles.

1.2 Distance functions

Def. $d: S \times S \rightarrow \mathbb{R}$ is *distance function* iff pos. definite except $d(x,x) = 0$ ($d(x,x') > 0 \iff x \neq x'$), symmetric ($d(x,x') = d(x',x)$) and triangle inequality holds ($d(x,x'') \leq d(x,x') + d(x',x'')$).

Euclidean L_r $d_r(x,y) = \|x - y\|_r = (\sum_i |x_i - y_i|^r)^{1/r}$.

Cosine $\operatorname{Sim}_c(A,B) = \frac{A \cdot B}{\|A\| \cdot \|B\|}$, $d_c(A,B) = \frac{\cos^{-1}(\operatorname{Sim}_c(A,B))}{\pi}$.

Jaccard sim., d. $\operatorname{Sim}_J(A,B) = \frac{|A \cap B|}{|A \cup B|}$, $d_J(A,B) = 1 - \operatorname{Sim}_J(A,B)$.

1.3 LSH – local sensitive hashing

Key Idea: Similar documents have similar hash.

Note: Trivial for exact duplicates (hash-collision \rightarrow candidate pair).

Min-hash-family $h_\pi(C)$ for Jaccard Hash is the *min* (i.e. first) non-zero permuted row index: $h_\pi(C) = \min_{i, C(i)=1} \pi(i)$, bin. vec. C , rand. perm. π .

Note: $\Pr_\pi[h_\pi(C_1) = h_\pi(C_2)] = \operatorname{Sim}_J(C_1, C_2)$ if $\pi \in \text{u.a.r. } S_{|C|}$.

Min-hash L_r -norm: Fix $a \in \mathbb{R}$. Random line \mathbf{w} partitioned in buckets of length a . Project \mathbf{x}, \mathbf{y} onto \mathbf{w} , if in same bucket, $h_{\mathbf{w}}(x) = h_{\mathbf{w}}(y)$. In 2-dim. forms a $(a/2, 2 \cdot a, 1/2, 1/3)$ -sensitive hash-family. In d-dim. there exists a $(d1, d2, p1, p2)$ -sensitive family $\forall d1 < d2$ with $p1 > p2$.

Min-hash cos. $h_{\mathbf{w}}(x) = \operatorname{sgn}(\mathbf{w}^T \mathbf{x})$. $\Pr_{\mathbf{w}}[h_{\mathbf{w}}(x) = h_{\mathbf{w}}(y)] = 1 - \frac{\theta_{\mathbf{w}, \mathbf{y}}}{\pi}$.

Min-hash signature matrix $M_S \in [N]^{n \times C}$ with $M_S(i, c) = h_i(C_c)$ given n hash-fns h_i drawn randomly from a universal hash family.

Pseudo permutation h_π with $\pi(i) = (a \cdot i + b) \bmod p \bmod N$, N number of shingles, $p \geq N$ prime and $a, b \in \text{u.a.r. } [p]$ with $a \neq 0$. Use as universal hash family. Only store a and b . Much more efficient.

Compute signature matrix M_S For column $c \in [C]$, row $r \in [N]$ with $C_c(r) = 1$, $M_S(i, c) \leftarrow \min\{h_i(C_c), M_S(i, c)\}$ for all h_i .

$(d1, d2, p1, p2)$ -sensitivity of $F = \{h_1, \dots, h_n\}$: $\forall x, y \in S: d(x, y) \leq d1 \implies P[h(x) = h(y)] \geq p1$ and $d(x, y) \geq d2 \implies P[h(x) = h(y)] \leq p2$.

r -way AND $h = [h_1, \dots, h_r]$, $h(x) = h(y) \iff \forall i \ h_i(x) = h_i(y)$ is $(d1, d2, p_1^r, p_2^r)$ -sensitive. Decreases FP.

b -way OR $h = [h_1, \dots, h_b]$, $h(x) = h(y) \iff \exists i \ h_i(x) = h_i(y)$ is $(d1, d2, 1 - (1 - p_1)^b, 1 - (1 - p_2)^b)$ -sensitive. Decreases FN.

Boosting by composition b -way OR after r -way AND. Group sig. matrix into b bands of r rows. CP match in at least one band (check by hashing). Result is $(d1, d2, 1 - (1 - p_1^r)^b, 1 - (1 - p_2^r)^b)$ -sensitive. r pulls stronger than b . r pulls down.

r -way AND after b -way OR. Result is $(d1, d2, (1 - (1 - p_1)^b)^r, (1 - (1 - p_2)^b)^r)$ -sensitive. b pulls stronger than r . b pulls up.

Tradeoff FP/FN Favor FP (work) over FN (wrong). Filter FP by checking signature matrix, shingles or even whole documents.

2 Supervised Learning

Linear classifier $y_i = \operatorname{sgn}(\mathbf{w}^T \mathbf{x}_i)$ assuming \mathbf{w} goes through origin.

Homogeneous transform $\tilde{\mathbf{x}} = [\mathbf{x}, 1]; \tilde{\mathbf{w}} = [\mathbf{w}, b]$, now \mathbf{w} passes origin.

Kernel k is inner product in high-dim. space: $k(\mathbf{x}, \mathbf{y}) = \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle$. *shift-invariance* $k(\mathbf{x}, \mathbf{y}) = k(\mathbf{x} - \mathbf{y})$.

Gaussian $k(\mathbf{x} - \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|_2^2 / h^2)$.

Convex function $f: S \rightarrow \mathbb{R}$ is convex iff $\forall x, x' \in S, \lambda \in [0, 1], \lambda f(x) + (1 - \lambda)f(x') \geq f(\lambda x + (1 - \lambda)x')$, i. e. every segment lies above function. Equiv. bounded by linear fn. at every point.

H-strongly convex f *H-strongly convex* iff $f(x') \geq f(x) + \nabla f(x)^T(x' - x) + \frac{H}{2}\|x' - x\|_2^2$, i. e. bounded by quadratic fn (at every point).

2.1 Support vector machine (SVM)

SVM primal

Quadratic $\min_{\mathbf{w}} \mathbf{w}^T \mathbf{w} + C \sum_i \xi_i$, s.t. $\forall i: y_i \mathbf{w}^T \mathbf{x}_i \geq 1 - \xi_i$, $\forall i: \xi_i \geq 0$.

Hinge loss $\min_{\mathbf{w}} \lambda \mathbf{w}^T \mathbf{w} + \sum_i \max(0, 1 - y_i \mathbf{w}^T \mathbf{x}_i)$ with $\lambda = \frac{1}{C}$.

Norm-constrained $\min_{\mathbf{w}} \sum_i \max(0, 1 - y_i \mathbf{w}^T \mathbf{x}_i)$ s.t. $\|\mathbf{w}\|_2 \leq \frac{1}{\sqrt{\lambda}}$.

Lagrangian dual $\max_{\alpha} \sum_i \alpha_i + \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$, $\alpha_i \in [0, C]$. Apply kernel trick: $\max_{\alpha} \sum_i \alpha_i + \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$, $\alpha_i \in [0, C]$, prediction becomes $y = \operatorname{sgn}(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x}))$.

2.2 Convex Programming

Convex program $\min_{\mathbf{x}} f(\mathbf{x})$, s. t. $\mathbf{x} \in S$, f convex.

Online convex program (OCP) $\min_{\mathbf{w}} \sum_{t=1}^T f_t(\mathbf{w})$, s. t. $\mathbf{w} \in S$.

General regularized form $\min_{\mathbf{w}} \sum_{i=1}^n l(\mathbf{w}; \mathbf{x}_i, y_i) + \lambda R(\mathbf{w})$, where l is a (convex) loss function and R is the (convex) regularizer.

General norm-constrained form $\min_{\mathbf{w}} \sum_{i=1}^n l(\mathbf{w}; \mathbf{x}_i, y_i)$, s. t. $\mathbf{w} \in S_\lambda$, l is loss and S_λ some (norm-)constraint. Note: This is an OCP.

Solving OCP *Feasible set* $S \subseteq \mathbb{R}^d$ and *start pt.* $\mathbf{w}_0 \in S$, OCP (as above). Round $t \in [T]$: pick feasible pt. \mathbf{w}_t , get convex fn. f_t , incur $l_t = f_t(\mathbf{w}_t)$. Regret $R_T = (\sum_{t=1}^T l_t) - \min_{\mathbf{w} \in S} \sum_{t=1}^T f_t(\mathbf{w})$.

Online SVM $\|\mathbf{w}\|_2 \leq \frac{1}{\lambda}$ (norm-constr.). For new pt. \mathbf{x}_t classify $y_t = \operatorname{sgn}(\mathbf{w}_t^T \mathbf{x}_t)$, incur $l_t = \max(0, 1 - y_t \mathbf{w}_t^T \mathbf{x}_t)$, update \mathbf{w}_t (see later). Best $L^* = \min_{\mathbf{w}} \sum_{t=1}^T \max(0, 1 - y_t \mathbf{w}^T \mathbf{x}_t)$, regret $R_t = \sum_{t=1}^T l_t - L^*$.

Online proj. gradient descent (OPGD) Update for online SVM: $\mathbf{w}_{t+1} = \operatorname{Proj}_S(\mathbf{w}_t - \eta_t \nabla f_t(\mathbf{w}_t))$ with $\operatorname{Proj}_S(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}' \in S} \|\mathbf{w}' - \mathbf{w}\|_2$, gives regret bound $\frac{R_T}{T} \leq \frac{1}{\sqrt{T}} (\|\mathbf{w}_0 - \mathbf{w}^*\|_2^2 + \|\nabla f\|_2^2)$, $\nabla f_t(\mathbf{w}_t) = -y_t \mathbf{x}_t$, if $y_t \mathbf{w}_t^T \mathbf{x}_t \leq 1$, else 0; $\eta_t = \frac{1}{\sqrt{t}}$ and $\operatorname{Proj}_S(\mathbf{w}') = \mathbf{w}' \cdot \min(1, 1/(\sqrt{\lambda} \|\mathbf{w}'\|_2))$.

For H -strongly convex fn, $\eta_t = \frac{1}{Ht}$ gives $R_t \leq \frac{\|\nabla f\|_2^2}{2H} (1 + \log T)$.

Stochastic Gradient Descent (SGD) Convex program is unconstrained and decomposes $f(\mathbf{w}) = \sum_i^n g(\mathbf{w}; i)$. In SVM $g(\mathbf{w}; i) = \frac{\lambda \mathbf{w}^T \mathbf{w}}{n} + l_h(\mathbf{w}; \mathbf{x}_i, y_i)$. *Algo:* $\mathbf{w} \leftarrow 0$ and choose β_t . until convergence: $(\mathbf{x}_i, y_i) \in \text{u.a.r. } X$. $\mathbf{w} \leftarrow \mathbf{w} - \beta_t \cdot \nabla_{\mathbf{w}}(g(\mathbf{w}; i))$.

Stochastic PGD (SPGD) Online-to-batch. Compute $\tilde{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}_t$. If data i. i. d.: exp. error (risk) $\mathbb{E}[L(\tilde{\mathbf{w}})] \leq \frac{L(\mathbf{w}^*)}{T} + R_T/T$, $L(\mathbf{w}^*)$ is best error (risk) possible.

PEGASOS OPGD w/ mini-batches on strongly convex SVM form. $\min_{\mathbf{w}} \sum_{t=1}^T g_t(\mathbf{w})$, s.t. $\|\mathbf{w}\|_2 \leq \frac{1}{\sqrt{t}}$, $g_t(\mathbf{w}) = \frac{\lambda}{2} \|\mathbf{w}\|_2^2 + f_t(\mathbf{w})$.

g_t is λ -strongly convex, $\nabla g_t(\mathbf{w}) = \lambda \mathbf{w} + \nabla f_t(\mathbf{w})$.

Performance ϵ -accurate sol. with prob. $\geq 1 - \delta$ in runtime $O^*(\frac{d \cdot \log \frac{1}{\delta}}{\lambda \epsilon})$.

ADAGrad Adapt to geometry. *Mahalanobis norm* $\|\mathbf{w}\|_G = \|\mathbf{G} \mathbf{w}\|_2$. $\mathbf{w}_{t+1} = \operatorname{argmin}_{\mathbf{w} \in S} \|\mathbf{w} - (\mathbf{w}_t - \eta \mathbf{G}_t^{-1} \nabla f_t(\mathbf{w}))\|_{\mathbf{G}_t}$. Min. regret with $G_t = (\sum_{\tau=1}^t \nabla f_\tau(\mathbf{w}_\tau) \nabla f_\tau(\mathbf{w}_\tau)^T)^{1/2}$. Easily inv'able matrix with $G_t = \operatorname{diag}(\dots)$. $R_t \in O(\frac{\|\mathbf{w}^*\|_\infty}{\sqrt{T}} \sqrt{d})$, even better for sparse data.

ADAM Add 'momentum' term: $\mathbf{w}_{t+1} = \mathbf{w}_t - \mu \bar{g}_t$, $g_t = \nabla f_t(\mathbf{w})$, $\bar{g}_t = (1 - \beta)g_t + \beta \bar{g}_{t-1}$, $\bar{g}_0 = 0$. Helps for dense gradients.

Parallel SGD (PSGD) Randomly partition to k (indep.) machines. Comp. $\mathbf{w} = \frac{1}{k} \sum_{i=1}^k \mathbf{w}_i$. $\mathbb{E}[\text{err}] \in O(\epsilon(\frac{1}{k\sqrt{\lambda}} + 1))$ if $T \in \Omega(\frac{\log \frac{k\lambda}{\epsilon}}{\epsilon \lambda})$. Suitable for MapReduce cluster, multi. passes possible.

Hogwild! Shared mem., no sync., sparse data. [...]

Implicit kernel trick Map $x \in \mathbb{R}^d \rightarrow \phi(x) \in \mathbb{R}^D \rightarrow z(x) \in \mathbb{R}^m$, $d \ll D, m \ll D$. Where $\phi(x)$ corresponds to a kernel $k(x, x') = \phi(x)^T \phi(x')$.

Random fourier features Given shift-invariant kernel k .

$p(\omega) = \frac{1}{2\pi} \int e^{-j\omega' \delta} k(\delta) d\Delta$
 $\omega_i \sim p = \text{eg Gaussian}$, $b_i \sim U(0, 2\pi)$

$z(\mathbf{x}) \equiv \sqrt{2/m} [\cos(\omega'_1 \mathbf{x} + b_1) \dots \cos(\omega'_m \mathbf{x} + b_m)]$

Nyström features (need entire dataset) In practice: pick random samples $S = \{\hat{\mathbf{x}}_1 \dots \hat{\mathbf{x}}_n\} \subseteq X$

$K_{S \times i, j} = k(\hat{\mathbf{x}}_i, \hat{\mathbf{x}}_j)$, $K_{SS, i, j} = k(\hat{\mathbf{x}}_i, \hat{\mathbf{x}}_j)$

approximate $\mathbf{K} = \mathbf{K}_{X \times S} \mathbf{K}_{SS}^{-1} \mathbf{K}_{S \times X}$, $\mathbf{K}_{SS} = \mathbf{V} \mathbf{D} \mathbf{V}^T$.

new point \mathbf{x}' : $z(\mathbf{x}') = \mathbf{D}^{-1/2} \mathbf{V}^T [k(\mathbf{x}', \hat{\mathbf{x}}_1), \dots, k(\mathbf{x}', \hat{\mathbf{x}}_m)]$

3 Pool-based active Learning (semi-supervised)

Uncertainty sampl. Until all labels can be inferred (or until a *label budget*); $U_t(x) = U(x|_{x_{1:t-1}, y_{1:t-1}})$, request y_t for $x_t = \operatorname{argmax}_x U_t(x)$. **SVM:** $x_t = \operatorname{argmin}_{x_i} |\mathbf{w}^T \mathbf{x}_i|$, i.e. $U_t(\mathbf{x}) = \frac{1}{|\mathbf{w}^T \mathbf{x}|}$.

Sub-linear time w/ LSH $|w^T x_i|$ small if \angle_{w,x_i} close to $\frac{\pi}{2}$ (90°). Hash hyperplane: $h_{u,v}(a,b) = [h_u(a), h_v(b)] = [\text{sgn}(u^T a), \text{sgn}(v^T b)]$. LSH hash family: $h_H(z) = h_{u,v}(z, z)$ if z datapoint, $h_H(z) = h_{u,v}(z, -z)$ if z query hyperplane. $\Pr[h_H(w) = h_H(x)] = \Pr[h_u(w) = h_u(x)]\Pr[h_v(-w) = h_v(x)] = \frac{1}{4} - \frac{1}{\pi^2}(\angle_{x,w} - \frac{\pi}{2})^2$. Hash all unlabeled. Loop: Hash w , req. labels for hash-coll., update.

Informativeness Metric of “information” gainable; \neq uncertainty.

Version Space $\mathcal{V}(D) = \{w \mid \forall (x,y) \in D \text{ sgn}(w^T x) = y\}$

Relevant version space given unlabeled pool $U = \{x'_1, \dots, x'_n\}$. $\tilde{\mathcal{V}}(D;U) = \{h: U \rightarrow \{\pm 1\} \mid \exists w \in \mathcal{V}(D) \forall x \in U \text{ sgn}(w^T x) = h(x)\}$.

Generalized binary search Init $D \leftarrow \{\}$. While $|\tilde{\mathcal{V}}(D;U)| > 1$, comp. $v^\pm(x) = |\tilde{\mathcal{V}}(D \cup \{(x, \pm)\}; U)|$, label of $\text{argmin}_x \max\{v^-(x), v^+(x)\}$.

Approx. $|\mathcal{V}|$ Margins of SVM $m^\pm(x)$ for labels $\{+, -\}, \forall x$. *Max-min* $\max_x \min\{m^+(x), m^-(x)\}$ or *ratio* $\max_x \min\{\frac{m^+(x)}{m^-(x)}, \frac{m^-(x)}{m^+(x)}\}$.

4 Model-based clustering – Unsupervised learning

k-means problem $\min_\mu L(\mu)$ with $L(\mu) = \sum_{i=1}^N \min_j \|\mathbf{x}_i - \mu_j\|_2^2$ and *cluster centers* $\mu = \mu_1, \dots, \mu_k$. Non-convex! NP-hard in general!

LLoyd's Init $\mu^{(0)}$ (somehow/randomly). *Assign* all \mathbf{x}_i to closest center $z_i \leftarrow \text{argmin}_{j \in [k]} \|\mathbf{x}_i - \mu_j^{(t-1)}\|_2^2$, *Update* to mean: $\mu_j^{(t)} \leftarrow \frac{1}{n_j} \sum_{i: z_i = j} \mathbf{x}_i$. Always converge to *local minimum*.

Online k-means Init μ somehow/randomly. For $t \in [n]$ find $c = \text{argmin}_j \|\mu_j - \mathbf{x}_t\|_2$, set $\mu_c \leftarrow \mu_c + \eta_t(\mathbf{x}_t - \mu_c)$. For local optimum: $\sum_t \eta_t = \infty \wedge \sum_t \eta_t^2 < \infty$ suffices, e.g. $\eta_t = \frac{c}{t}$, $c \in \mathbb{R}$.

Weighted rep. C $L_k(\mu; C) = \sum_{(w,x) \in C} w \cdot \min_j \|\mu_j - x\|_2^2$.

(k, ϵ) -coreset iff $\forall \mu: (1 - \epsilon)L_k(\mu; D) \leq L_k(\mu; C) \leq (1 + \epsilon)L_k(\mu; D)$.

D^2 -sampling Sample prob. $p(x) = \frac{d(x,B)^2}{\sum_{x' \in X} d(x',B)^2}$.

Merge coresets union of (k, ϵ) -coreset is also (k, ϵ) -coreset.

Compress a (k, δ) -coreset of a (k, ϵ) -coreset is a $(k, \epsilon + \delta + \epsilon\delta)$ -coreset.

Coresets on streams Bin. tree of merge-compress. Error \propto height.

Mapreduce k-means Construct (k, ϵ) -coreset C , solve k-means (w/ many restarts) on coreset. (Repeat.) Near-optimal solution.

5 k -armed bandits as recommender systems

k -armed bandit k arms with diff. prob. dist. For $t \in [T]$ rounds, pick $i_t \in [k]$, sample $y_t \in P_i$ (indep. of other rounds). Max. $\sum_{t=1}^T y_t$.

Regret μ_i mean of P_i (arm i), $\mu^* = \max_i \mu_i$. Regret $r_t = \mu^* - \mu_{i_t}$. Total regret $R_T = \sum_{t=1}^T r_t$.

ϵ -greedy *Explore* u.a.r. with prob. ϵ_t , *exploit* with prob. $1 - \epsilon_t$: choose $\text{argmax}_i \hat{\mu}_i$. Suitable $\epsilon_t \in O(1/t)$ gives $R_T \in O(k \log T)$. Clearly unoptimal.

UCB1 Init $\hat{\mu}_i \leftarrow 0$; try all arms once. Following $t \in [T - k]$ rounds: $UCB(i) \leftarrow \hat{\mu}_i + \sqrt{\frac{2 \log t}{n_i}}$, pick $i_t \leftarrow \text{argmax}_i UCB(i)$, observe y_t . Update $n_{i_t} \leftarrow n_{i_t} + 1, \hat{\mu}_{i_t} \leftarrow \hat{\mu}_{i_t} + \frac{y_t - \hat{\mu}_{i_t}}{n_{i_t}}$.

Contextual bandits Round t : Obs. *context* $\mathbf{z}_t \in \mathcal{Z} \subseteq \mathbb{R}^d$; *recommend* $\mathbf{x}_t \in \mathcal{A}_t$. Reward $y_t = f(\mathbf{x}_t, \mathbf{z}_t) + \epsilon_t$. Regret $r_t = \max_{\mathbf{x}} f(\mathbf{x}, \mathbf{z}_t) - f(\mathbf{x}_t, \mathbf{z}_t)$. Often $f(\mathbf{x}, \mathbf{z}) = \mathbf{w}_{\mathbf{x}}^T \mathbf{z}$ linear.

Idea behind LinUCB Estimate $\hat{\mathbf{w}}_i = \text{argmin}_{\mathbf{w}} \sum_{t=1}^m (y_t - \mathbf{w}^T \mathbf{z}_t) + \|\mathbf{w}\|_2^2$. Closed form: $\hat{\mathbf{w}}_i = M_i^{-1} D_i^T \mathbf{y}_i$, $M_i = D_i^T D_i + I$, $D_i = [\mathbf{z}_1 | \dots | \mathbf{z}_m], \mathbf{y}_i = (y_1 | \dots | y_m)^T$.

Confidence If $\alpha = 1 + \sqrt{\ln(\frac{2}{\delta})/2}$:

$\Pr \left[|\hat{\mathbf{w}}_i^T \mathbf{z}_t - \mathbf{w}_i^T \mathbf{z}_t| \leq \alpha \sqrt{\mathbf{z}_t^T M_i^{-1} \mathbf{z}_t} \right] \geq 1 - \delta$.

LinUCB (Algorithm) For $t = [T]$ receive *action set* A_t and features \mathbf{z}_t . For all $\mathbf{x} \in A_t$: if \mathbf{x} new, set $M_{\mathbf{x}} \leftarrow \mathbb{I}$ and $b_{\mathbf{x}} \leftarrow 0$; set $\hat{\mathbf{w}}_{\mathbf{x}} \leftarrow M_{\mathbf{x}}^{-1} b_{\mathbf{x}}$; set $UCB_{\mathbf{x}} \leftarrow \hat{\mathbf{w}}_{\mathbf{x}}^T \mathbf{z}_t + \alpha \sqrt{\mathbf{z}_t^T M_{\mathbf{x}}^{-1} \mathbf{z}_t}$. Recommend action $\mathbf{x}_t = \text{argmax}_{\mathbf{x} \in A_t} UCB_{\mathbf{x}}$; observe y_t . Set $M_{\mathbf{x}} \leftarrow M_{\mathbf{x}} + \mathbf{z}_t \mathbf{z}_t^T$ and $b_{\mathbf{x}} \leftarrow b_{\mathbf{x}} + y_t \mathbf{z}_t$.

Hybrid Model $y_t = \mathbf{w}_i^T \mathbf{z}_t + \beta^T \phi(\mathbf{x}_i, \mathbf{z}_t) + \epsilon_t$ captures sep. and shared effects.

Rejection Sampling Evaluate bandit: For $t \in \mathbb{N}$ read $\log(\mathbf{x}_1^{(t)}, \dots, \mathbf{x}_k^{(t)}, \mathbf{z}_t, a_t, y_t)$. Pick a'_t by algo. If $a'_t = a_t$ feed y_t to algo., else ignore line. Stop after T feedbacks.

6 Submodularity

$F: 2^V \rightarrow \mathbb{R}$ *subm.* iff $F(A \cup \{s\}) - F(A) \geq F(B \cup \{s\}) - F(B)$, $\forall A \subseteq B \subseteq V, s \in V$ with $s \notin B$.

Closedness of non-negative linear combination: If $F_1, \dots, m(A)$ subm. then $\lambda_i \geq 0: F'(A) := \sum_i \lambda_i F_i(A)$ subm.

Other closedness: If $F(S)$ is subm. on V and $W \subseteq V$, then *Restriction* $F(S \cap W)$, *Conditioning* $F(S \cup W)$, *Reflection* $F(V \setminus S)$ are submodular.

Marginal gain: $\Delta_F(s|A) = F(\{s\} \cup A) - F(A)$

Greedy algo: In round $i+1$, previously picked $A_i = \{s_1, \dots, s_i\}$; pick $s_{i+1} = \text{argmax}_s \Delta_F(s|A_i) = \text{argmax}_s (F(\{s\} \cup A_i) - F(A_i))$. This gives a $(1 - 1/e) \approx 0.63$ -approximation.

Lazy Greedy: *Observation:* Submodularity implies $\Delta(s|A_i) \geq \Delta(s|A_{i+1})$. Algo.: $A_0 \leftarrow \{\}$; first iteration as usual. Then keep ordered list of Δ_i from prev. iteration. For $i \in [k]$ do: $\Delta_i = F(A_{i-1} \cup \{s^*\})$, $s^* = \text{argmax}_s \Delta_F(s|A_{i-1})$ (*top element*). If s^* is still top, then $A_i \leftarrow A_{i-1} \cup \{s^*\}$ else resort and pick top element (as in Greedy).

7 Tips

$\sum \|y_i - w^T x_i\|^2 = \sum (y_i - w^T x_i)^T (y_i - w^T x_i) = (y - Xw)^T (y - Xw)$

Showing active learning needs n labels Assume points are all distinct and algo has returned -1 for the first $n-1$ labels. Choose x s.t. algo can not distinguish between $-1, +1$.

Sublinear strategy approach Sort and rename points. Cleverly search witnesses (points determining interval).

Deriving Dual Start with quadratic. Rewrite constraints $c_{1,i}, c_{2,i}$ s.t. $c_i(w) \geq 0$. Get Lagrangian $L(w, \xi, \alpha, \gamma)$: $f(w) - \sum_i \alpha_i c_{1,i} - \sum_i \gamma_i c_{2,i}$. Set derivatives $\frac{\partial L}{\partial w}$ and $\frac{\partial L}{\partial \xi_i}$ to 0 and get w, C . Insert into $L(w)$. Result is max. qp with constraint $\alpha_i \in [0, C]$

8 Probability

$$\begin{aligned} \mathbb{E}[X] &= \sum_x \Pr(x) \cdot x \\ \mathbb{E}[aX + aY] &= a \cdot \mathbb{E}[x] + b \cdot \mathbb{E}[y] \\ \text{Var}[X] &= \sum_x p(x) \cdot (x - \mathbb{E}[X])^2 \\ &= \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 \\ \text{Cov}[X, Y] &= \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])] = \mathbb{E}[X \cdot Y] - \mathbb{E}[X] \cdot \mathbb{E}[Y] \\ \text{Var}[\alpha X + \beta Y] &= \alpha \text{Var}[X] + 2\alpha\beta \text{Cov}[X, Y] + \beta \text{Cov}[Y] \\ \text{Cov}[X, X] &= \text{Var}[X] \\ \text{Cov}[A+B, X] &= \text{Cov}[A, X] + \text{Cov}[B, X] \\ \Pr(A \cup B) &= \Pr(A) + P(B) - \Pr(A \cap B) \\ \Pr(A \cup B) &= \Pr(A) + P(B) \text{ if A, B mutually exclusive} \\ \Pr(A \cap B) &= \Pr(A|B)P(B) = \Pr(B|A)P(A) \\ \Pr(A \cap B) &= \Pr(A)P(B) \text{ if A, B independent} \\ \Pr(A|B) &= \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)} \end{aligned}$$

Hoeffding's Inequality

$$\begin{aligned} \Pr(\overline{X} - \mathbb{E}[\overline{X}] \geq t) &\leq \exp\left(-\frac{2n^2 t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \\ \Pr(|\overline{X} - \mathbb{E}[\overline{X}]| \geq t) &\leq 2 \exp\left(-\frac{2n^2 t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \end{aligned}$$

Markov's Inequality $\mathbb{P}(X \geq a) \leq \frac{\mathbb{E}(X)}{a}$.

Gaussian $\mathcal{N}(\mu, \sigma^2)$: $\frac{1}{\sqrt{2\sigma^2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$